

Detector Ranking is Not Architecture-Invariant under Medical Covariate Shift

Anonymous CVPR submission

Paper ID ****

Abstract

Post-hoc out-of-distribution (OOD) detectors are almost always ranked on a single backbone. **Holding a medical covariate split fixed, logit-based scores reorder substantially with the backbone while feature-space scores (Mahalanobis, k NN) stay near the top throughout.** On Camelyon17 hospital-2 (WILDS), Kendall's $W = 0.694$ across eight CNN families and seven scores (seed 42). MSP AUROC spans 0.515 (ResNet-50) to 0.883 (DenseNet-121)—a 0.37 swing, as large as Mahalanobis—MSP on that ResNet. Copying ResNet-50's best logit score (ReAct) onto ConvNeXt costs 0.185 AUROC versus ConvNeXt's own best logit. On the same labeled OOD test set, the AUROC-highest backbone has coverage@risk 10% = 0.078 while the coverage@risk-highest backbone reaches 0.904: AUROC is not a reliable proxy for coverage@risk. Feature-space scores occupy ranks 1–2 on every official backbone. Skin ISIC→PAD does not reproduce the MSP jump ($W = 0.791$). AUROC versus coverage@risk disagreement flips sign between Camelyon and MIDOG. A pre-registered ViT-B/16 MSP of 0.884 met a locked threshold of 0.6947. The protocol: do not copy an OOD ranking across backbones, and do not pick a backbone for triage from an AUROC table.

1. Introduction

We test one hypothesis: *under a fixed medical covariate shift, post-hoc detector ranking is an architecture property.* On Camelyon17 hospital-2, ResNet-50 MSP AUROC is 0.515 while DenseNet-121 is 0.883 (Fig. 1a). A lab that copies an OpenMIBOOD-style ResNet-50 recipe and a lab that uses DenseNet are not ranking the same detectors; and, measured on the same labeled test set, AUROC order does not predict coverage@risk order, so an AUROC table is not a safe guide to triage performance—the AUROC-best and coverage@risk-best backbones differ by 0.826 coverage at a 10% risk budget (Fig. 1b).

That hypothesis is *not* that ranking was previously assumed architecture-invariant. On natural-image semantic OOD, ranking already moves with architecture [4, 6, 26, 34]. OpenMIBOOD [11] already showed that ImageNet method ranking does not transfer to medical data and that feature-space scores outperform logit scores under a frozen classifier per domain. The orthogonal question is: **holding one medical covariate-shift dataset fixed, how much does ranking move when only the backbone changes?**

The rest of the paper is that test, in order. Camelyon discovers the shift: Kendall $W = 0.694$ and a backbone-dependent MSP jump. Five seeds verify the DenseNet MSP jump (5/5). Matched-recipe training asks whether that jump is architecture or optimizer. Skin falsifies a universal MSP-jump law, while feature ranks stay 1–2. MIDOG is the deployment readout: the *sign* of AUROC versus coverage@risk disagreement flips. A pre-registered ViT-B/16 tests whether the Camelyon MSP jump extends past CNNs. The contribution is a protocol, not a detector.

Contributions.

- Camelyon shows architecture reshuffles the logit-score ranking.** Official Kendall $W = 0.694$ ($n=8$, 7 methods, seed 42). MSP AUROC span 0.515–0.883 matches the Maha—MSP gap on ResNet-50. Copying ResNet-50's winning logit (ReAct) onto ConvNeXt costs 0.185 AUROC. Feature-space scores stay in ranks 1–2. The seed-robust MSP jump is DenseNet-121 (0.265 ± 0.065 , 5/5); ConvNeXt and EffV2-S@224 are 4/5, MobileNet/RegNet 3/5, EffB3 2/5.
- Skin falsifies universality.** ISIC→PAD ($n=8$, $W=0.791$): Maha rank-1 and k NN rank-2 on all eight; MSP stays in a narrow band. Not a second Camelyon jump.
- AUROC is not a reliable proxy for deployment coverage.** The sign of the AUROC versus coverage@risk 10% disagreement flips between Camelyon and MIDOG [9, 29]. On the Camelyon $n=8$ zoo, measured on the same labeled test set, the AUROC-best and coverage@risk-best backbones differ by 0.826 coverage at a 10% risk




Figure 1. **The claim in one figure.** (a) Same Camelyon hospital-2 split, same seven scores: switching ResNet-50→DenseNet-121 lifts logit AUROC by ~ 0.37 while Mahalanobis and k NN stay on top. (b) On the same labeled test set, the AUROC-highest backbone (DenseNet) has coverage@risk 10% = 0.078 while the coverage@risk-highest backbone (MobileNet) reaches 0.904: AUROC rank and coverage rank disagree.

Work	#arch.	Shift	Rank stat.	Utility
OpenOOD v1 [32]	1/set	semantic	—	AUROC
OpenOOD v1.5 [34]	3	semantic	qualitative	AUROC
Szyc <i>et al.</i> [26]	3 var.	CIFAR sem.	rank Δ	AUROC
Datko <i>et al.</i> [6]	per line	natural	claimed stable	AUROC
OpenMIBOOD [11]	1/domain	medical	—	AUROC
Ours	8+ViT	covar.	Kendall W	AUROC+cov.

Table 1. Closest ranking-vs-backbone papers. Architecture-sensitive ranking exists on natural-image semantic OOD; the increment here is a held-fixed medical covariate split, zoo-scale W , and operational coverage.

budget; this is a retrospective comparison, not a validated selection rule (Sec. 4.5).

4. **A locked prediction.** Before training: ViT-B/16 Camelyon MSP ≥ 0.6947 . Observed 0.884. Held out of the 8-CNN W .

2. Related work

OpenMIBOOD [11] ranks 24 post-hoc detectors with a **single frozen classifier per medical domain**. It shows that ImageNet method ranking does not transfer to medical data and that feature-space scores outperform logit scores under that protocol. Table 1 locates the increment: we hold the dataset fixed and change the backbone.

OpenOOD v1 [32] fixes architecture on purpose. OpenOOD v1.5 [34] compares three ImageNet models on semantic near/far OOD and reports that “different

post-processor may favor different architecture,” without Kendall’s W and without a covariate hospital split. Szyc *et al.* [26] show rankings change across three ResNet CIFAR implementation variants. Datko *et al.* [6] claim an architecture-specific winner that transfers across natural-image OOD—the opposite of a single-split concordance statistic.¹ Claros Olivares and Brockmeier [4] rank confidence scores with AURC/AUGRC on CIFAR / TinyImageNet; they do not report a sign-flip of AUROC versus coverage@risk 10% across medical covariate domains.

OoD-Bench [33] is OOD generalization, not detection. WILDS [17] supplies Camelyon17 and iWildCam; iWildCam is a negative bound, not a third main domain.

3. Setup

3.1. Shifts

Camelyon17 [2], via the WILDS [17] split: hospitals 0, 3, 4 as ID and hospital 2 as OOD (WILDS’s official OOD-test domain); founding cell, official W , $n=8$. Hospital 1 is WILDS’s official OOD-validation domain; we do not use it for any number in the main paper and report it only in the supplement’s held-out stain-jitter check. Skin ISIC→PAD [5, 22]: feature rank-stability without an MSP jump. MIDOG [1]: OpenMIBOOD’s [11] 1a vs. 1b+1c

¹This citation could not be independently verified against a primary source at the time of writing; see `refs.bib`. Confirm the original DOI/URL before submission.

split and checkpoints, coverage sign-flip; official W is $n=3$. We build on OpenMIBOOD’s public MIDOG split and ResNet-50 checkpoint as an external reference. Public R50 AUROC ~ 0.59 is *theirs*; in-house R50 MSP 0.512 is *ours*. We never mix the two.

3.2. Backbones

Eight CNN-family models, ImageNet initialization, 10-epoch fine-tune, seed 42, train-frac 1.0 on Camelyon: ResNet-18/50 [12], DenseNet-121 [15], ConvNeXt-Tiny [21], MobileNetV3-Large [14], RegNetY-3.2GF [23], EfficientNet-B3 [27] at native 300, and EfficientNetV2-S [28] at 224 (native 384 dropped). Standard fine-tune recipes (Adam / AdamW / SGD) follow each family’s common practice and are *not* forced-identical Adam. That con-found is measured in Sec. 4.3. Only resolution was isolated in a controlled pair (EffB3 @224 vs. @300; EffV2-S @224 vs. @384). ViT-B/16 [7] @224 is held out of W .

3.3. Scores and statistics

Seven post-hoc scores: MSP [13], Energy [20], energy of logit-normed logits (ELogitNorm), ViM [30], ReAct [24], Mahalanobis with Ledoit–Wolf covariance [18, 19], and k NN [25]. Primary concordance: **Kendall’s W per domain, never pooled**. Jump: $\text{MSP} - \text{mean}(\text{R18}, \text{R50}) \geq 0.15$. Feature drag: $\text{Maha or } k\text{NN on a new backbone below median}(\text{R18}, \text{R50}, \text{EffB3}) - 0.10$. Official Camelyon ResNet mean $\text{MSP} = 0.544654$; jump iff $\text{MSP} \geq 0.695$ on that split.

Coverage@risk 10% is OOD-triage: rank OOD samples by MSP (higher = more ID-like) and take the largest coverage whose classifier risk is $\leq 10\%$. ID-selective Camelyon is vacuous (ID-val acc $\approx 99.5\%$) and is not cited. AUROC and coverage@risk measure different things: AUROC is ID-vs-OOD separability of the detector score alone, while coverage@risk 10% depends jointly on that ranking *and* the classifier’s own accuracy on the OOD domain, so the two need not agree, and a disagreement between them is not itself surprising without separately checking which one moved. Both quantities in Sec. 4.5 are computed directly on the labeled OOD test set; that makes the AUROC-vs-coverage@risk comparison there retrospective, not an executable backbone-selection procedure (Limitation 7).

We tested whether the Camelyon MSP pattern is an artifact of input resolution, temperature scaling [10], ID-only detector *selection*, a CE+SupCon recipe [16], or train-time HED stain jitter. None is a clean account (Sec. 5 and the supplement).

4. Results

4.1. Camelyon-17 discovers the ranking shift

Official Kendall $W = 0.694$ ($n=8$, 7 methods). This is the only Camelyon number called “official W .” Leave-one-backbone recomputation of W ranges 0.676–0.738; that is *structural* sensitivity of the concordance statistic, not a 95% CI (supplement).

Feature methods occupy ranks 1–2 on every official backbone (Maha #1 on 6/8; k NN #1 on EffB3 and MobileNet). No backbone drags Maha/ k NN past the 0.10 feature-drag threshold. On Table 2, MSP spans 0.515–0.883; Maha 0.870–0.988; k NN 0.856–0.973. Incomplete concordance ($W=0.694$) is the logit ranks moving (Fig. 2).

Table 3 reports the full five-seed MSP matrix: seeds verify the DenseNet MSP jump. DenseNet-121 is the confirmed seed-robust jump (5/5). ConvNeXt-Tiny and EfficientNetV2-S@224 are not seed-42 flukes (4/5) and are not promoted to 5/5. MobileNetV3 and RegNetY are 3/5; EfficientNet-B3 is 2/5 (its seed-42 jump $+0.258$ is a favorable draw relative to 0.142 ± 0.125). The seed-42 matrix still shows a jump without feature drag. The claim is narrower than “6/6 non-ResNet, confirmed.”

Copying a logit ranking is not free. ResNet-50’s best logit is ReAct (0.722); on ConvNeXt the same score is 0.692 against Energy 0.877—a 0.185 AUROC transfer cost. Feature-space transfer regret is ≤ 0.05 because Maha/ k NN stay in ranks 1–2. Architecture moves the logit leaderboard; it does not unseat the feature default.

4.2. Skin falsifies a universal MSP jump

Official $W = 0.791$ ($n=8$). Maha is rank-1 and k NN rank-2 on **all eight** (Fig. 2). MSP sits in a narrow band (0.68–0.74). There is no Camelyon-style ResNet-vs-rest MSP jump. EffV2-S@224 MSP is 0.737; the native-384 cell that put k NN at rank 7 (MSP 0.823) is a resolution artifact and is not in official W . Skin is rank-stability of feature scores under a zoo: a replication of OpenMIBOOD’s feature>logit pattern *across architectures*, not an MSP-jump.

4.3. Optimizer challenges a clean architecture reading

Resolution is not the Camelyon DenseNet / seed-42 pattern. EffB3 Camelyon: 0.803 @300 $\rightarrow$ 0.787 @224 (the jump versus the official ResNet mean survives). EffV2-S Camelyon: 0.726 @384 $\rightarrow$ **0.778** @224 (the official cell; the jump is not a 384 artifact). On skin, the EffV2-S @384 k NN collapse *is* a resolution artifact and disappears at 224. These two facts are not a slogan that “resolution never matters.”

M1 — jumpers trained with ResNet Adam ($\text{lr}=10^{-4}$, $\text{wd}=10^{-4}$, batch 64). The four official zoo members

Backbone	input	Maha (rank)	kNN (rank)	MSP	seed-42 jump
ResNet-18	224	0.956 (1)	0.936 (2)	0.574	no
ResNet-50	224	0.896 (1)	0.866 (2)	0.515	no
DenseNet-121	224	0.988 (1)	0.973 (2)	0.883	yes
ConvNeXt-Tiny	224	0.928 (1)	0.916 (2)	0.846	yes
MobileNetV3-L	224	0.870 (2)	0.904 (1)	0.771	yes
RegNetY-3.2GF	224	0.964 (1)	0.956 (2)	0.784	yes
EfficientNet-B3	300	0.885 (2)	0.934 (1)	0.803	yes
EfficientNetV2-S	224	0.880 (1)	0.856 (2)	0.778	yes

Table 2. Camelyon hospital-2 AUROC, **seed 42** (official W). Jump iff $\text{MSP} \geq 0.6947$ vs. frozen ResNet mean 0.544654. MSP range is ResNet-50→DenseNet.

Backbone	42	43	44	45	46	jump mean $\pm$ sd	≥ 0.15
ResNet-18	0.574	0.712	0.601	0.524	0.581	—	—
ResNet-50	0.515	0.621	0.619	0.515	0.661	—	—
ResNet mean	0.545	0.667	0.610	0.520	0.621	—	—
DenseNet-121	0.883	0.865	0.868	0.845	0.826	0.265 $\pm$ 0.065	5/5
ConvNeXt-Tiny	0.846	0.799	0.803	0.764	0.785	0.207 $\pm$ 0.067	4/5
MobileNetV3-L	0.771	0.728	0.815	0.780	0.745	0.175 $\pm$ 0.081	3/5
RegNetY-3.2GF	0.784	0.778	0.836	0.830	0.724	0.198 $\pm$ 0.089	3/5
EfficientNet-B3	0.803	0.731	0.643	0.818	0.678	0.142 $\pm$ 0.125	2/5
EfficientNetV2-S@224	0.778	0.744	0.808	0.846	0.835	0.210 $\pm$ 0.089	4/5

Table 3. Same eight Camelyon hospital-2 MSP cells on all five seeds {42, 43, 44, 45, 46}. Jump = MSP minus that seed’s ResNet mean; bar 0.15. The last column is how many of those five seeds clear the bar. Official W is seed 42 only.

whose standard recipe was not that Adam still jump versus the frozen official ResNet mean 0.545: ConvNeXt 0.820 (+0.275), MobileNetV3 0.737 (+0.193), RegNet 0.805 (+0.261), EffV2-S@224 0.873 (+0.328); all with id-val acc ≥ 0.994 . 4/4 clean hits. DenseNet and EffB3 already shared ResNet-Adam in the official zoo. Recipe mismatch does *not* explain the seed-42 jumps of those four members. M1 is seed 42 only; the five-seed MSP matrix is Table 3.

M2 — ResNets trained with ConvNeXt AdamW (wd=0.05). ResNet-18 MSP **0.773** (enters the 0.695 band; id-val 0.9957). ResNet-50 MSP 0.581 (stays out; id-val 0.9958). ResNet *can* jump under a non-default recipe. Architecture is not a clean variable, independent of optimizer. These two results point in opposite directions; we report both.

4.4. ViT-B/16: a locked prediction

Locked *before* training, from the seed-42 zoo rule that non-ResNet-like cells jumped: $\text{MSP} \geq 0.6947$. Linear depth/parameter/receptive-field predictors are invalid on this $n=8$ and are not quoted. Observed MSP **0.884**, Maha 0.958, k NN 0.941. The prediction holds. Not a ninth W column.

4.5. Coverage@risk: the deployment cost

Protocol: OOD-triage coverage@risk 10% with MSP, not ID-selective coverage.

On Camelyon hospital-2, the locked $n=4$ utility table has MSP AUROC order EffB3@300 (0.803) > EffB3@224 (0.787) > ResNet-18 (0.574) > ResNet-50 (0.515), but coverage@risk 10% order ResNet-50 (0.149) > ResNet-18 (0.145) > EffB3@300 (0.107) > EffB3@224 (0.013). The AUROC winner is not the coverage winner (Fig. 3).

On MIDOG, official $n=3$ (in-house): MSP AUROC ResNet-50 0.512, EffB3 0.486, ResNet-18 0.438 (gaps are small). Coverage@risk 10%: EffB3 0.712, ResNet-18 0.442, ResNet-50 0.166. ResNet-50 is AUROC-highest and coverage-lowest. The permutation is *not* the Camelyon permutation.

AURC/AUGRC already assume that AUROC and coverage-based ranking can disagree [9, 29]. The increment is the *sign* of that disagreement flipping between two medical covariate domains under the same protocol, with no rule for which backbone wins. Both quantities above are computed on the same labeled Camelyon OOD test set: this is a retrospective comparison of two different measurements of the same eight backbones, not an executable backbone-selection rule, since coverage@risk requires OOD labels that are not available at deployment time. What it shows




Figure 2. **Method rank by backbone** (1=best AUROC). Feature-space scores occupy ranks 1–2 on every official backbone in both domains. Logit ranks move on Camelyon ($W=0.694$) and much less on skin ($W=0.791$), where Maha is rank-1 on 8/8.




Figure 3. **AUROC vs. coverage@risk 10%**. On Camelyon the AUROC-high EfficientNets have the *lowest* triage coverage; on MIDOG the AUROC-high ResNet-50 is coverage-lowest. The sign of the disagreement flips.

is that AUROC order and coverage@risk order diverge by 0.826 coverage at a 10% risk budget on this zoo (DenseNet 0.078 vs. MobileNet 0.904; Fig. 1b), so AUROC is not a reliable proxy for coverage@risk. A deployable estimate of coverage@risk should instead come from a held-out validation split of the target domain, not from a source-domain AUROC table; we flag this as an open protocol gap (Limitation 7) rather than claim a selection method here.

4.6. Frozen foundation models

Camelyon linear-probe MSP does not invert: GigaPath [31] 0.664, Phikon [8] 0.605, UNI [3] 0.599, all >0.5 . MIDOG UNI / GigaPath MSP 0.445 / 0.470 is a MIDOG-only

footnote. FM Maha ~ 0.995 –1.0 is discussion support that feature>logit can survive a training-paradigm change, not a pillar.

5. Failed alternatives

Each of the following was a precommitted test. None is hidden.

Resolution (from Sec. 4.3). EffB3@224 MSP 0.787; EffV2-S@224 MSP 0.778. Not a pixel-size explanation of Table 2.

Temperature T^* . Correlation of T^* with the MSP jump: $r=0.268$, $p=0.52$. DenseNet–ResNet-18 MSP gap

0.308→0.308 after T^* . Calibration/temperature is not the mechanism.

ID-only Maha/MSP gate. Selecting a score from ID-only statistics is a *selector*, not a detector. On Camelyon it picks Maha on 2/8 backbones; mean AUROC of the gated score 0.837 versus always-Maha 0.940. The bar was matching Maha and ViM. It does not.

CE+SupCon ($\lambda=0.1$, $\tau=0.07$), ResNet-18 + DenseNet. Gap 0.308→0.238 (23%; precommit bar was 50% / ≤ 0.154). ResNet-18 MSP 0.574→0.515 (worse; this is the SupCon cell, not official ResNet-50 MSP 0.515). Part of the shrink is the low model degrading. Dirty; not a recipe.

HED stain jitter, an in-house MIDOG $n=8$ zoo, and a shift-type extract (covariate / task / far-OOD) are in the supplement. Joint readout: no shift-type law, no stain mitigation, no Camelyon-style MSP jump on MIDOG.

6. Discussion and limitations

The hypothesis holds on Camelyon and fails as a cross-medical MSP-jump law. Architecture reshuffles the logit-score ranking substantially—ResNet versus the rest—while feature-space scores (Mahalanobis, k NN) stay near the top across every backbone we tested; $W=0.694$ is itself a fairly high concordance, and the two score families are not equally sensitive to the backbone, so “architecture governs ranking” is true of the logit scores and not of the feature-space scores. Feature-space scores remain the stable default when the shift looks like Camelyon *and* when it does not (skin). M1 preserves the seed-42 jumps under ResNet-Adam; M2 lets ResNet-18 enter the MSP band under ConvNeXt-AdamW: architecture is the dominant axis, not a clean one. ViT-B/16 was a locked prediction, not a post-hoc ninth CNN.

The practical cost is not abstract. On the labeled Camelyon test set, the AUROC-best and coverage@risk-best backbones differ by 0.83 of OOD coverage at a 10% risk budget; whether that gap is still there when coverage@risk is estimated from a held-out validation split rather than read off the test labels is open (Limitation 7). Copying ResNet-50’s logit winner onto ConvNeXt costs 0.185 AUROC. Those are two different kinds of transfer error—one from reading a selection off an AUROC table, one from reusing a logit score across backbones—not from picking a weak detector in isolation.

Limitations. (1) **iWildCam bound.** Locked 3-backbone cell: DenseNet MSP 0.575 vs. ResNet mean 0.604, jump -0.029 . Extra five CNNs vs. that domain’s ResNet mean: ConvNeXt 0.601, MobileNet 0.558, RegNet 0.580, EffB3 0.577, EffV2-S 0.650; 0/5 jump. The Camelyon MSP-jump is absent on this non-medical covariate shift

and may be biomedical-specific. iWildCam is not pooled with Camelyon W . (2) **No clean mitigation.** CE+SupCon is dirty (ResNet-18 MSP worse). HED stain jitter is dirty (DenseNet MSP $0.780 < 0.80$). The ID-only gate is a selector. (3) **MIDOG official W is $n=3$** ($W=0.857$); the in-house $n=8$ zoo shows no MSP jump and does not replace it. (4) W is per-domain; 0.694 / 0.791 / 0.857 are not averaged. (5) No reader study: the operational claim stops at coverage@risk. Coverage@risk itself is a point estimate here; per-sample bootstrap CIs are left to the toolkit release (features are not re-extracted in this draft). (6) Architecture is confounded with its standard recipe. 5/5 seed-robustness is DenseNet only. Official W stays seed 42. Staining vs. scanner vs. tissue prep inside the hospital shift is not isolated. (7) **The coverage@risk backbone comparison (Sec. 4.5) is retrospective.** Both the AUROC ranking and the coverage@risk ranking there are computed on the same labeled Camelyon OOD test set; this is a diagnostic comparison of two measurements on the same backbones, not a procedure that can be run without OOD labels. A proper selection protocol would estimate coverage@risk on a held-out validation split of the target domain and report it on a disjoint test split; that experiment is specified but not yet run (decision_precommit_coverage_val_split.md). (8) **No cross-seed concordance control.** Official $W=0.694$ uses seed 42 only, and Table 3 shows ResNet-18 MSP alone ranges from 0.524 to 0.712 across five seeds. We have not computed Kendall’s W across seeds of one fixed architecture with all seven scores, which would bound how much of $W=0.694$ ’s incomplete concordance could instead be seed variance; the existing five-seed data logs MSP only and cannot answer this. That experiment is specified but not yet run (decision_precommit_cross_seed_w.md).

Code for the architecture-zoo evaluation and coverage@risk will be released, crediting OpenMIBOOD for MIDOG imglists, the public ResNet-50 checkpoint, and the 1a / cs-ID / near-OOD protocol.

7. Conclusion

Under medical covariate shift, logit-based detector scores reorder substantially with the backbone while feature-space scores stay near the top throughout: architecture governs part of the ranking, but it does not override the feature-space family’s advantage. On Camelyon that is a seed-robust DenseNet MSP jump, stable feature ranks, and a 0.83-coverage gap between the AUROC-best and coverage@risk-best backbone, both measured on the same labeled test set. Skin keeps the feature ranks and drops the MSP jump. MIDOG flips the sign of the AUROC–coverage disagreement. A pre-registered ViT hit shows the Camelyon logit instability is not a CNN-zoo accident; iWildCam does not reproduce it. The protocol is the result: rank detectors

on the backbone you will deploy, and estimate that backbone's coverage@risk on a held-out validation split of the deployment shift, not from a source-domain AUROC table.

References

- [1] Marc Aubreville, Nikolas Stathonikos, Christof A. Bertram, Robert Klopffleisch, Natalie ter Hoeve, Francesco Ciompi, Frauke Wilm, Christian Marzahl, Taryn A. Donovan, Andreas Maier, Jack Breen, Nishant Ravikumar, Youjin Chung, Jinah Park, Ramin Nateghi, Fattaneh Pourakpour, Rutger H. J. Fick, Saima Ben Hadj, Mostafa Jahanifar, Adam Shephard, and Katharina Breininger. Mitosis domain generalization in histopathology images — the MIDOG challenge. *Medical Image Analysis*, 84:102699, 2023.
- [2] Péter Bándi, Oscar Geessink, Quirine Manson, Marcory Van Dijk, Maschenka Balkenhol, Meyke Hermesen, Babak Ehteshami Bejnordi, Byungjae Lee, Kyunghyun Paeng, Aoxiao Zhong, Quanzheng Li, Farhad Ghazvinian Zanjani, Svitlana Zinger, Keisuke Fukuta, Daisuke Komura, Vlado Ovtcharov, Shenghua Cheng, Shaoqun Zeng, Jeppe Thagaard, Anders B. Dahl, Huangjing Lin, Hao Chen, Ludwig Jacobsson, Martin Hedlund, Melih Çetin, Eren Halici, Hunter Jackson, Richard Chen, Fabian Both, Jörg Franke, Heidi Küsters-Vandeveld, Willem Vreuls, Peter Bult, Bram van Ginneken, Jeroen van der Laak, and Geert Litjens. From detection of individual metastases to classification of lymph node status at the patient level: The CAMELYON17 challenge. *IEEE Transactions on Medical Imaging*, 38(2):550–560, 2019.
- [3] Richard J. Chen, Tong Ding, Ming Y. Lu, Drew F. K. Williamson, Guillaume Jaume, Andrew H. Song, Bowen Chen, Andrew Zhang, Daniel Shao, Muhammad Shaban, et al. Towards a general-purpose foundation model for computational pathology. *Nature Medicine*, 30:850–862, 2024. 5
- [4] Esteban Claros Olivares and Austin J. Brockmeier. Confidence score ranking under distribution shift. *arXiv:2511.11934*, 2025. 1, 2
- [5] Noel Codella, Veronica Rotemberg, Philipp Tschandl, M. Emre Celebi, Stephen Dusza, David Gutman, Brian Helba, Aadi Kalloo, Konstantinos Liopyris, Michael Marchetti, Harald Kittler, and Allan Halpern. Skin lesion analysis toward melanoma detection 2018: A challenge hosted by the international skin imaging collaboration (ISIC). *arXiv:1902.03368*, 2019.
- [6] Szymon Datko, Kamil Szyc, Tomasz Walkowiak, and Henryk Maciejewski. Architecture-specific post-hoc out-of-distribution detection. *Expert Systems with Applications*, 2026. 1, 2
- [7] Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, and Neil Houlsby. An image is worth 16x16 words: Transformers for image recognition at scale. In *ICLR*, 2021. 3
- [8] Alexandre Filiot, Ridouane Ghermi, Antoine Olivier, Paul Jacob, Lucas Fidon, Alice Camara, Axel Mac Kain, Charlie Saillard, and Jean-Baptiste Schiratti. Scaling self-supervised learning for histopathology with masked image modeling. *medRxiv*, 2023. 5
- [9] Yonatan Geifman and Ran El-Yaniv. SelectiveNet: A deep neural network with an integrated reject option. In *ICML*, 2019. 1, 4
- [10] Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *ICML*, 2017. 3
- [11] Max Gutbrod, David Rauber, Danilo Weber Nunes, and Christoph Palm. OpenMIBOOD: Open medical imaging benchmarks for out-of-distribution detection. In *CVPR*, pages 25874–25886, 2025. 1, 2
- [12] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *CVPR*, 2016. 3
- [13] Dan Hendrycks and Kevin Gimpel. A baseline for detecting misclassified and out-of-distribution examples in neural networks. In *ICLR*, 2017. 3
- [14] Andrew Howard, Mark Sandler, Grace Chu, Liang-Chieh Chen, Bo Chen, Mingxing Tan, Weijun Wang, Yukun Zhu, Ruoming Pang, Vijay Vasudevan, et al. Searching for MobileNetV3. In *ICCV*, 2019. 3
- [15] Gao Huang, Zhuang Liu, Laurens van der Maaten, and Kilian Q. Weinberger. Densely connected convolutional networks. In *CVPR*, 2017. 3
- [16] Prannay Khosla, Piotr Teterwak, Chen Wang, Aaron Sarna, Yonglong Tian, Phillip Isola, Aaron Maschinot, Ce Liu, and Dilip Krishnan. Supervised contrastive learning. In *NeurIPS*, 2020. 3
- [17] Pang Wei Koh, Shiori Sagawa, Henrik Marklund, Sang Michael Xie, Marvin Zhang, Akshay Balsubramani, Weihua Hu, Michihiro Yasunaga, Richard Lanus Phillips, Irena Gao, et al. WILDS: A benchmark of in-the-wild distribution shifts. In *ICML*, 2021. 2
- [18] Olivier Ledoit and Michael Wolf. A well-conditioned estimator for large-dimensional covariance matrices. *Journal of Multivariate Analysis*, 88(2):365–411, 2004. 3
- [19] Kimin Lee, Kibok Lee, Honglak Lee, and Jinwoo Shin. A simple unified framework for detecting out-of-distribution samples and adversarial attacks. In *NeurIPS*, 2018. 3
- [20] Weitang Liu, Xiaoyun Wang, John Owens, and Yixuan Li. Energy-based out-of-distribution detection. In *NeurIPS*, 2020. 3
- [21] Zhuang Liu, Hanzi Mao, Chao-Yuan Wu, Christoph Feichtenhofer, Trevor Darrell, and Saining Xie. A ConvNet for the 2020s. In *CVPR*, 2022. 3
- [22] Andre G. C. Pacheco, Gustavo R. Lima, Amanda S. Salomão, Breno Krohling, Igor P. Biral, Gabriel G. de Angelo, Fábio C. R. Alves Jr, José G. M. Esgario, Alana C. Simora, Pedro B. C. Castro, et al. PAD-UFES-20: A skin lesion dataset composed of patient data and clinical images collected from smartphones. *Data in Brief*, 32:106221, 2020. 2
- [23] Ilija Radosavovic, Raj Prateek Kosaraju, Ross Girshick, Kaiming He, and Piotr Dollár. Designing network design spaces. In *CVPR*, 2020. 3

- 493 [24] Yiyu Sun, Chuan Guo, and Yixuan Li. ReAct: Out-of-
494 distribution detection with rectified activations. In *NeurIPS*,
495 2021. 3
- 496 [25] Yiyu Sun, Yifei Ming, Xiaojin Zhu, and Yixuan Li. Out-of-
497 distribution detection with deep nearest neighbors. In *ICML*,
498 2022. 3
- 499 [26] Kamil Szyc, Tomasz Walkowiak, and Henryk Maciejewski.
500 The reliability of OOD detection is in the details. In *UAI*,
501 2023. 1, 2
- 502 [27] Mingxing Tan and Quoc V. Le. EfficientNet: Rethinking
503 model scaling for convolutional neural networks. In *ICML*,
504 2019. 3
- 505 [28] Mingxing Tan and Quoc V. Le. EfficientNetV2: Smaller
506 models and faster training. In *ICML*, 2021. 3
- 507 [29] Jeremias Traub, Till J. Bungert, Constantin T. Lüth, Michael
508 Baumgartner, Klaus H. Maier-Hein, and Paul F. Jaeger.
509 Overcoming common flaws in the evaluation of selective
510 classification systems. In *NeurIPS*, 2024. 1, 4
- 511 [30] Haoqi Wang, Zhizhong Li, Litong Feng, and Wayne Zhang.
512 ViM: Out-of-distribution with virtual-logit matching. In
513 *CVPR*, 2022. 3
- 514 [31] Hanwen Xu, Naoto Usuyama, Jaspreet Bagga, Sheng Zhang,
515 Rajesh Rao, Tristan Naumann, Cliff Wong, Zelalem Gero,
516 Javier González, Yu Gu, et al. A whole-slide foundation
517 model for digital pathology from real-world data. *Nature*,
518 630:181–188, 2024. 5
- 519 [32] Jingkan Yang, Pengyun Wang, Dejian Zou, Zitang Zhou,
520 Kaiwen Ding, Wenbo Peng, Haoqi Wang, Guangyao Chen,
521 Bo Li, Yiyu Sun, et al. OpenOOD: Benchmarking general-
522 ized out-of-distribution detection. In *NeurIPS Datasets and*
523 *Benchmarks*, 2022. 2
- 524 [33] Nanyang Ye, Kaican Li, Haoyue Bai, Runpeng Yu, Lanqing
525 Hong, Fengwei Zhou, Zhenguo Li, and Jun Zhu. OoD-
526 Bench: Quantifying and understanding two dimensions of
527 out-of-distribution generalization. In *CVPR*, 2022. 2
- 528 [34] Jingyang Zhang, Jingkan Yang, Pengyun Wang, Haoqi
529 Wang, Yueqian Lin, Haoran Zhang, Yiyu Sun, Xuefeng
530 Liu, et al. OpenOOD v1.5: Enhanced benchmark for out-
531 of-distribution detection. *arXiv:2306.09301*, 2023. 1, 2